



CSCI-GA.3033-012

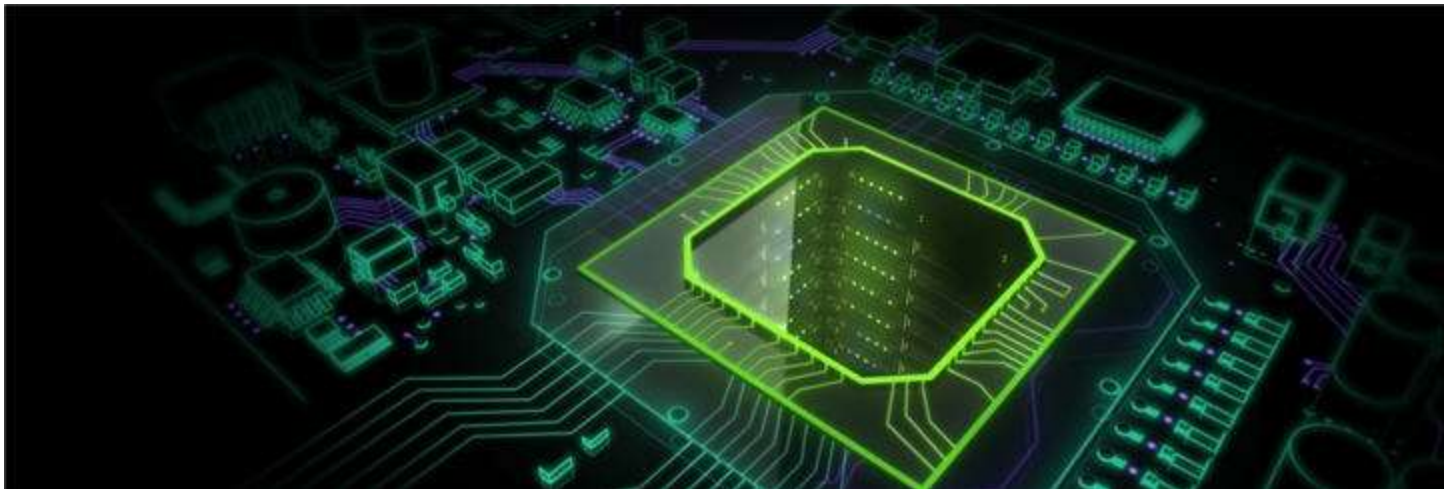
Graphics Processing Units (GPUs): Architecture and Programming

Lecture 1: Introduction

Mohamed Zahran (aka Z)

mzahran@cs.nyu.edu

<http://www.mzahran.com>



Who Am I?



- Mohamed Zahran (aka Z)
- Computer architecture/OS/Compilers Interaction
- <http://www.mzahran.com>
- Office hours: Wed 5:00-7:00 pm
- Room: WWH 328
- Course web page:

<http://cs.nyu.edu/courses/spring12/CSCI-GA.3033-012/index.html>

Formal Goals of This Course

- Why GPUs
- GPU Architecture
- GPU-CPU Interaction
- GPU programming model
- Solving real-life problems using GPUs

Informal Goals of This Course

- To get more than an A
- To learn GPUs and enjoy it
- To use what you have learned in *MANY* different contexts
- To have a feeling about how hardware and software evolve

The Course Web Page


- Lecture slides
- Info about mailing list, labs,
- Useful links (manuals, tools, book errata, ...)

The Textbook

Programming Massively Parallel
Processors: A Hands-on Approach

By


David B. Kirk & Wen-mei W. Hwu



Grading

- Homework assignments : 10%
- Project : 20%
- Programming assignments : 30%
- Final : 40%

Computer History




ENIAC

Eckert and Mauchly



- 1st working electronic computer (1946)
- 18,000 Vacuum tubes
- 1,800 instructions/sec
- 3,000 ft³

Computer History



EDSAC 1 (1949)

<http://www.cl.cam.ac.uk/UoCCL/misc/EDSAC99/>

- Maurice Wilkes




1st stored program
computer

650 instructions/sec

1,400 ft³


Intel 4004 Die Photo



- Introduced in 1970
 - First microprocessor
- 2,250 transistors
- 12 mm²
- 108 KHz




Intel 8086 Die Scan



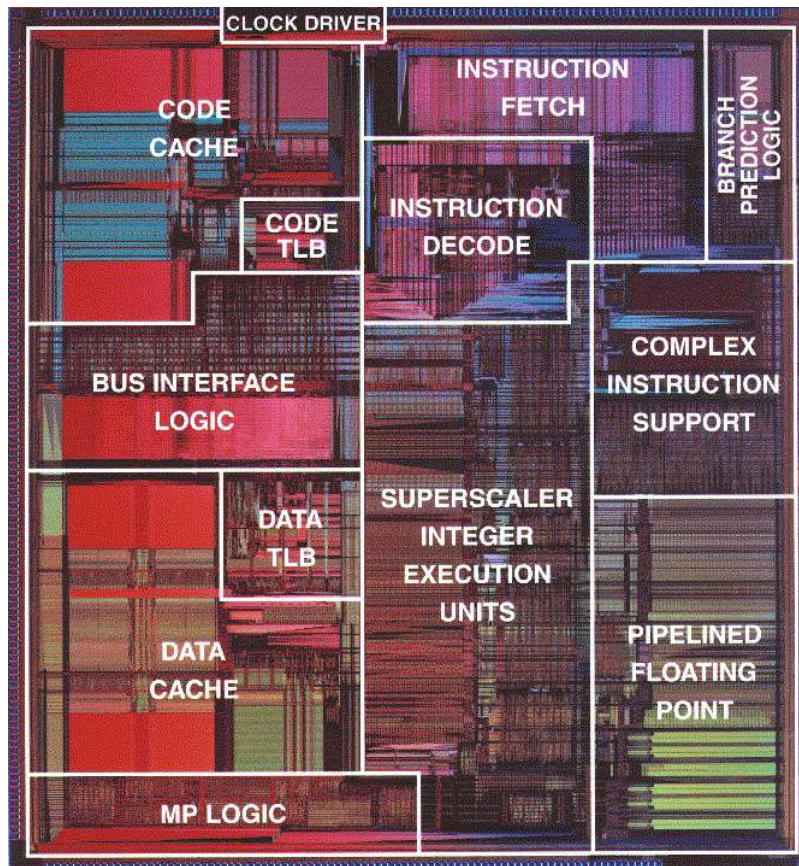
- 29,000 transistors
- 33 mm²
- 5 MHz
- Introduced in 1979
 - Basic architecture of the IA32 PC

Intel 80486 Die Scan




- 1,200,000 transistors
- 81 mm²
- 25 MHz
- Introduced in 1989
 - 1st pipelined implementation of IA32

Pentium Die Photo



- 3,100,000 transistors
- 296 mm²
- 60 MHz
- Introduced in 1993
 - 1st superscalar implementation of IA32

Pentium III



- 9,500,000 transistors
- 125 mm²
- 450 MHz
- Introduced in 1999


http://www.intel.com/intel/museum/25anniv/hof/hof_main.htm

Pentium 4




- 55,000,000 transistors
- 146 mm²
- 3 GHz
- Introduced in 2000


<http://www.chip-architect.com>




Pentium 4




Core 2 Duo (Merom)




Intel Core i7 (Nehalem)




IBM Power 7



Montecito




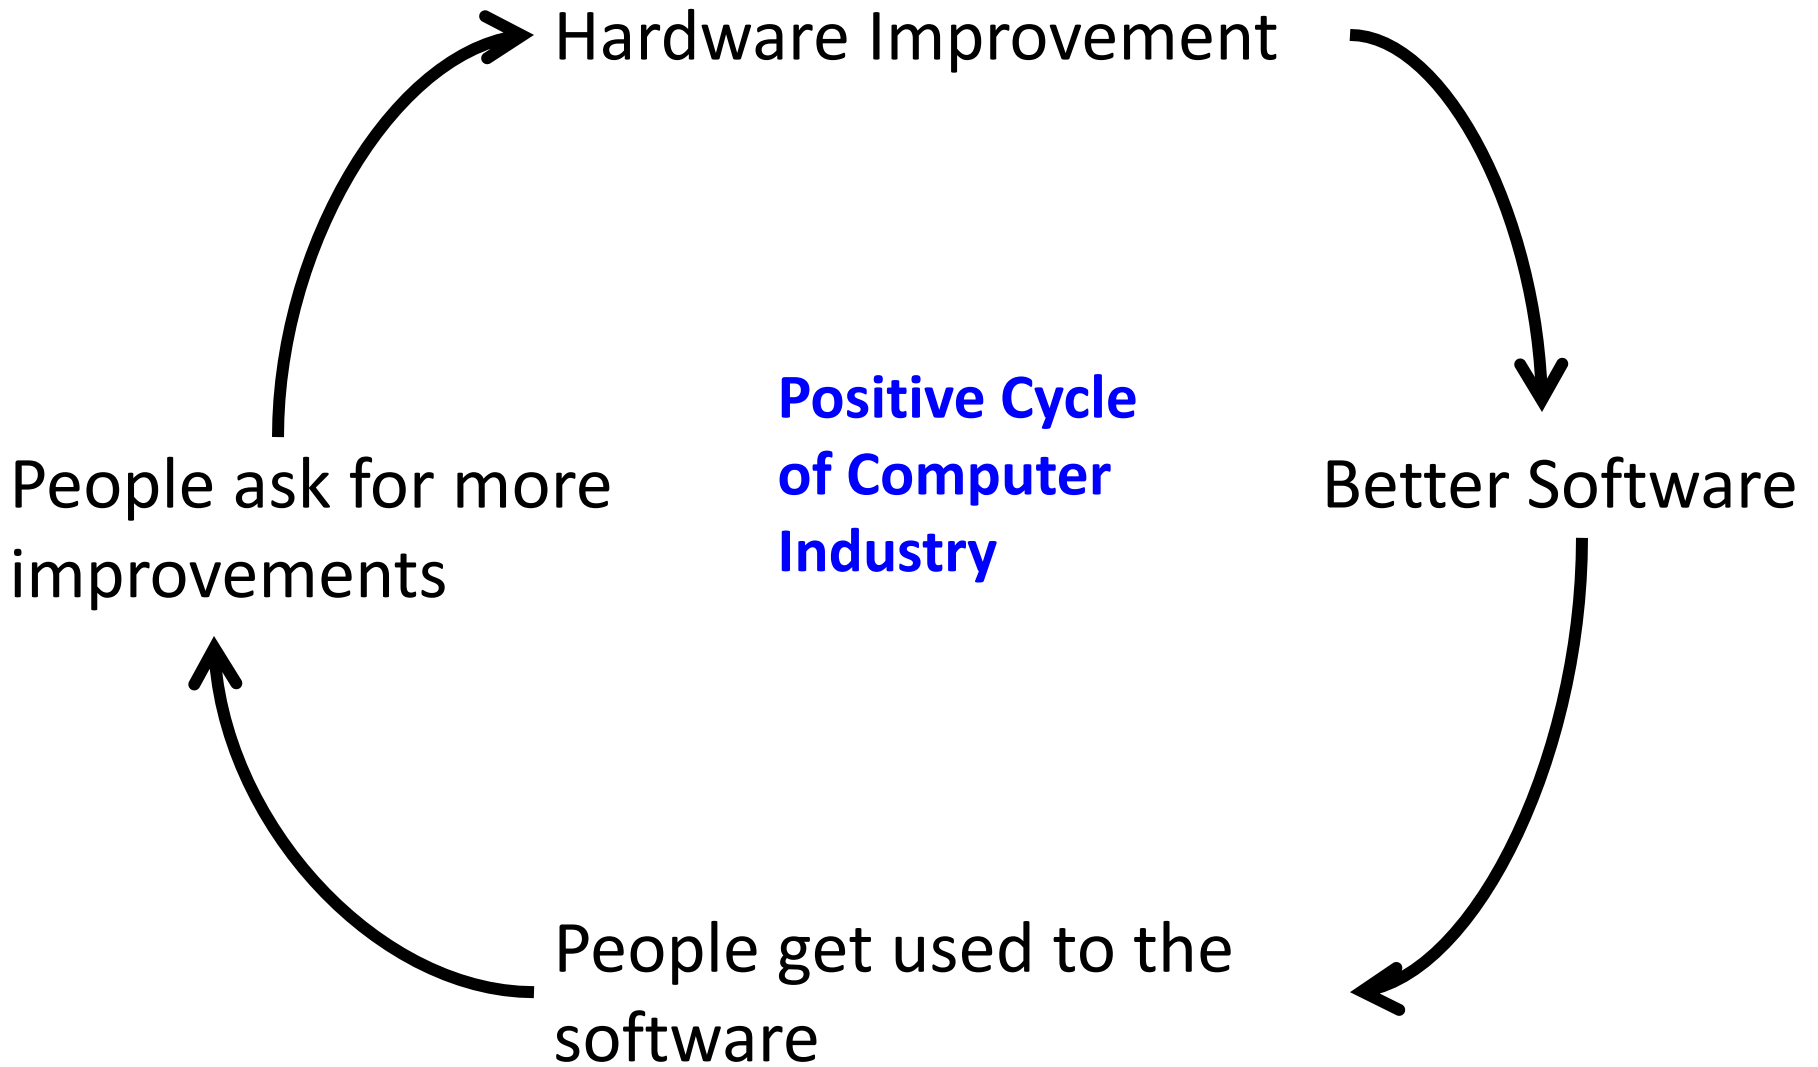
Cell Processor



Niagara
(SUN UltraSparc T2)

The Famous Moore's Law





The Status-Quo

- We moved from single core to multicore
 - for technological reasons
- Free lunch is over for software folks
 - The software will not become faster with every new generation of processors
- Not enough experience in parallel programming
 - Parallel programs of old days were restricted to some elite applications -> very few programmers
 - Now we need parallel programs for many different applications

Two Main Goals

- Maintain execution speed of old sequential programs
- Increase throughput of parallel programs

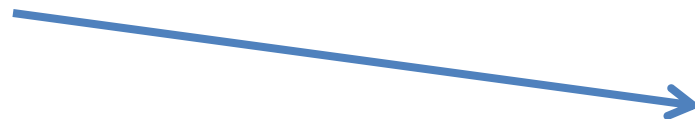
Two Main Goals

- Maintain execution speed of old sequential programs

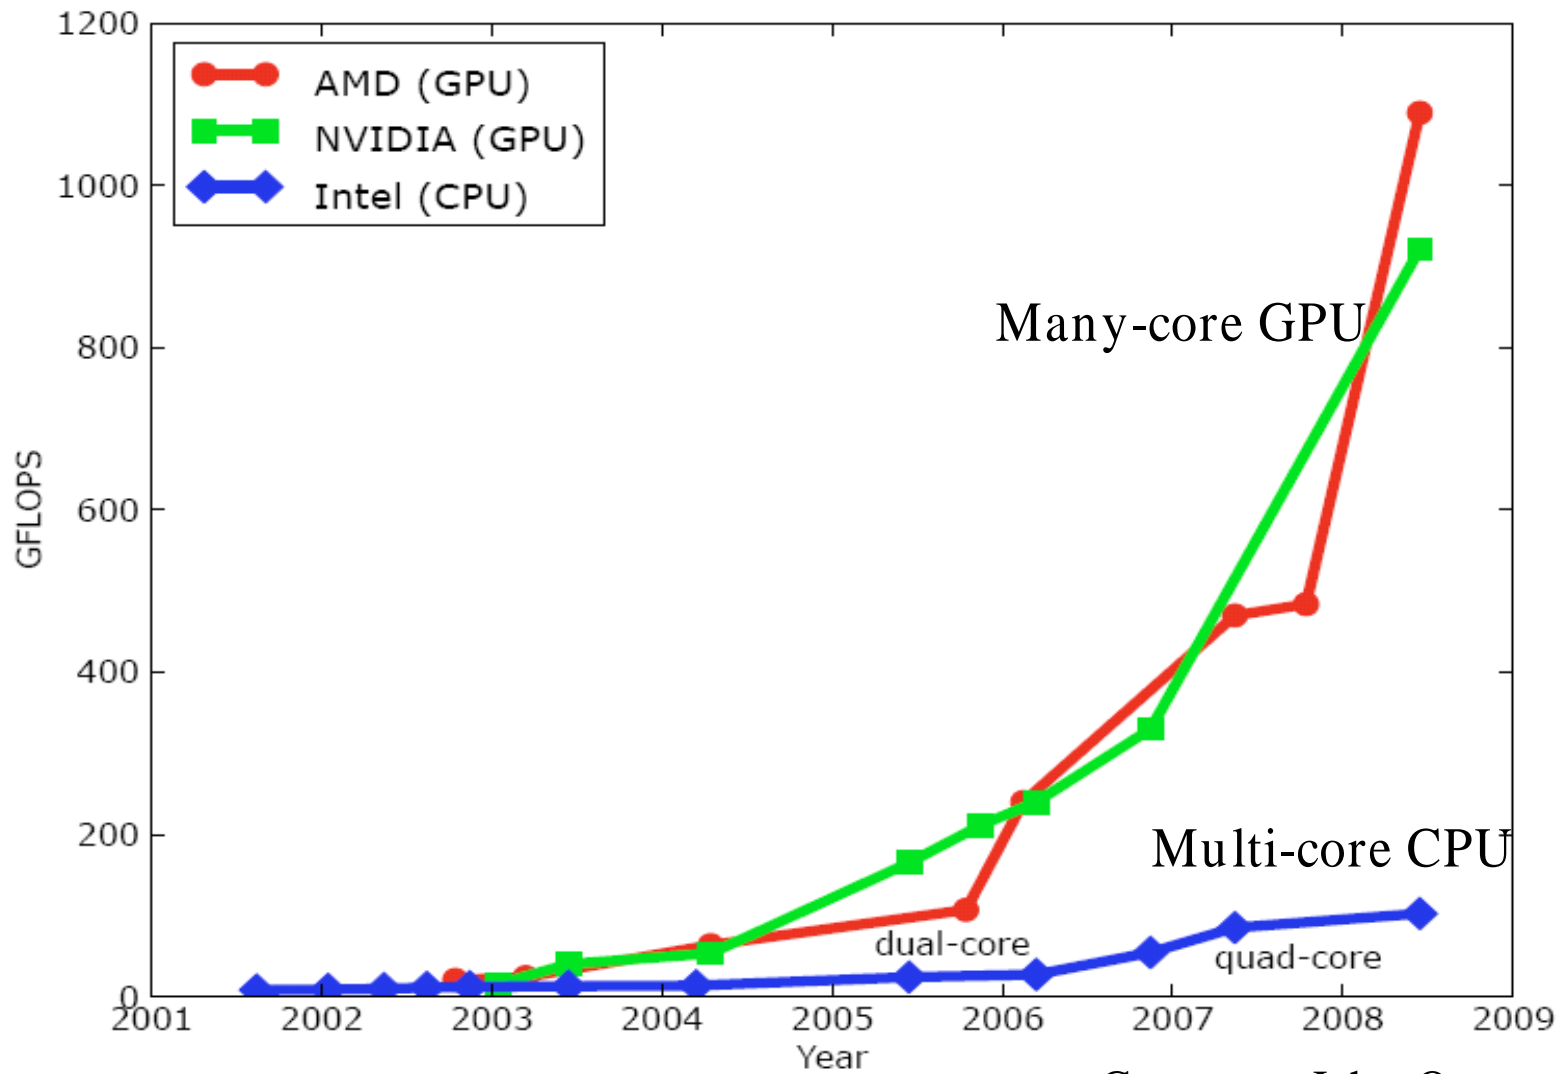


CPU

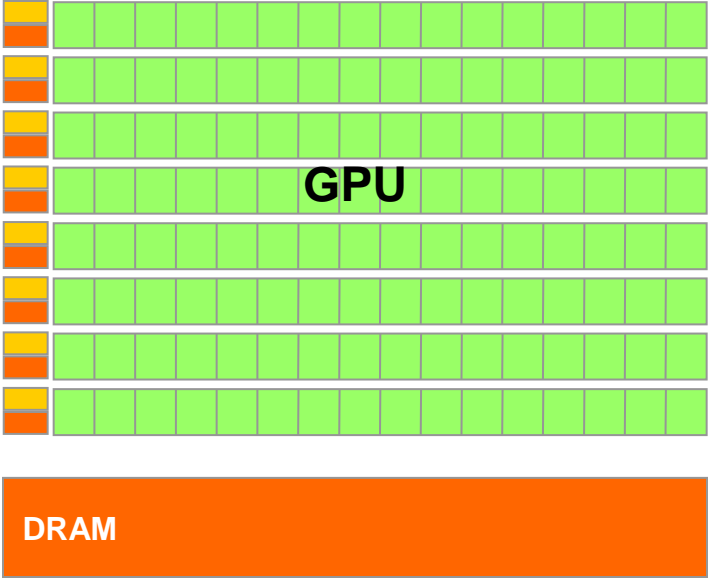

- Increase throughput of parallel programs




GPU





Courtesy: John Owens



**CPU is optimized for sequential
code performance**





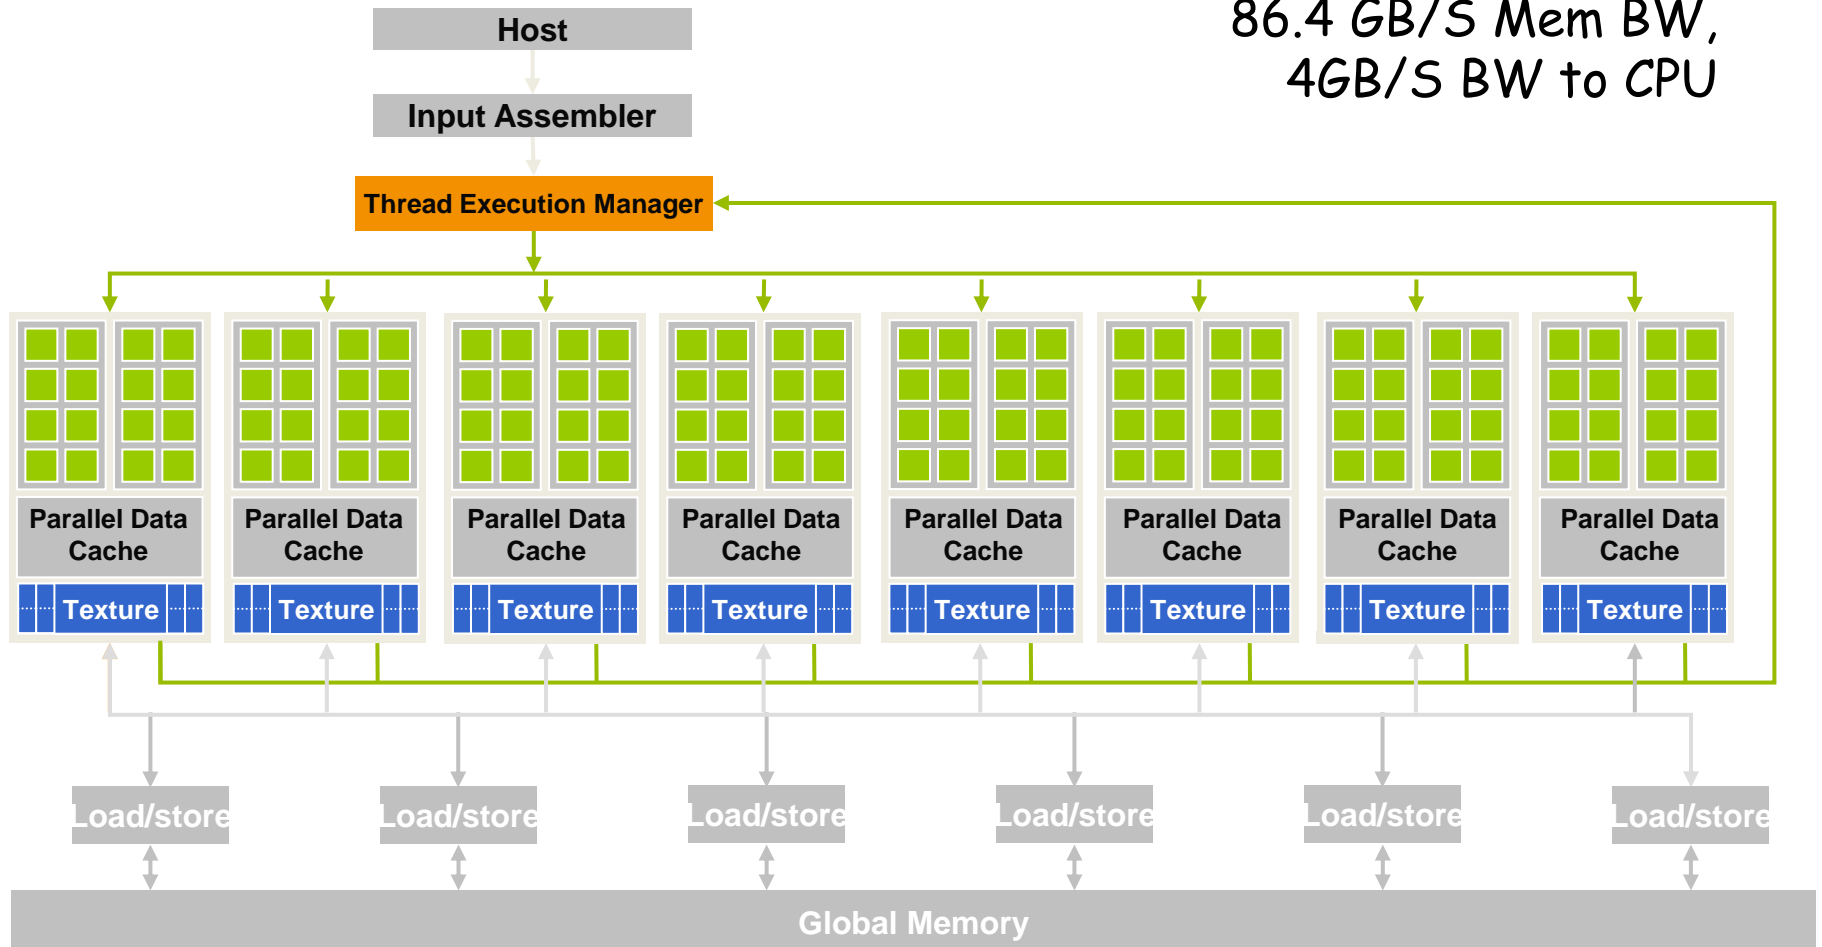
Almost 10x the bandwidth of multicore
(relaxed memory model)

How to Choose A Processor for Your Application?

- Performance
- Very large installation base
- Practical form-factor and easy accessibility
- Support for IEEE floating point standard

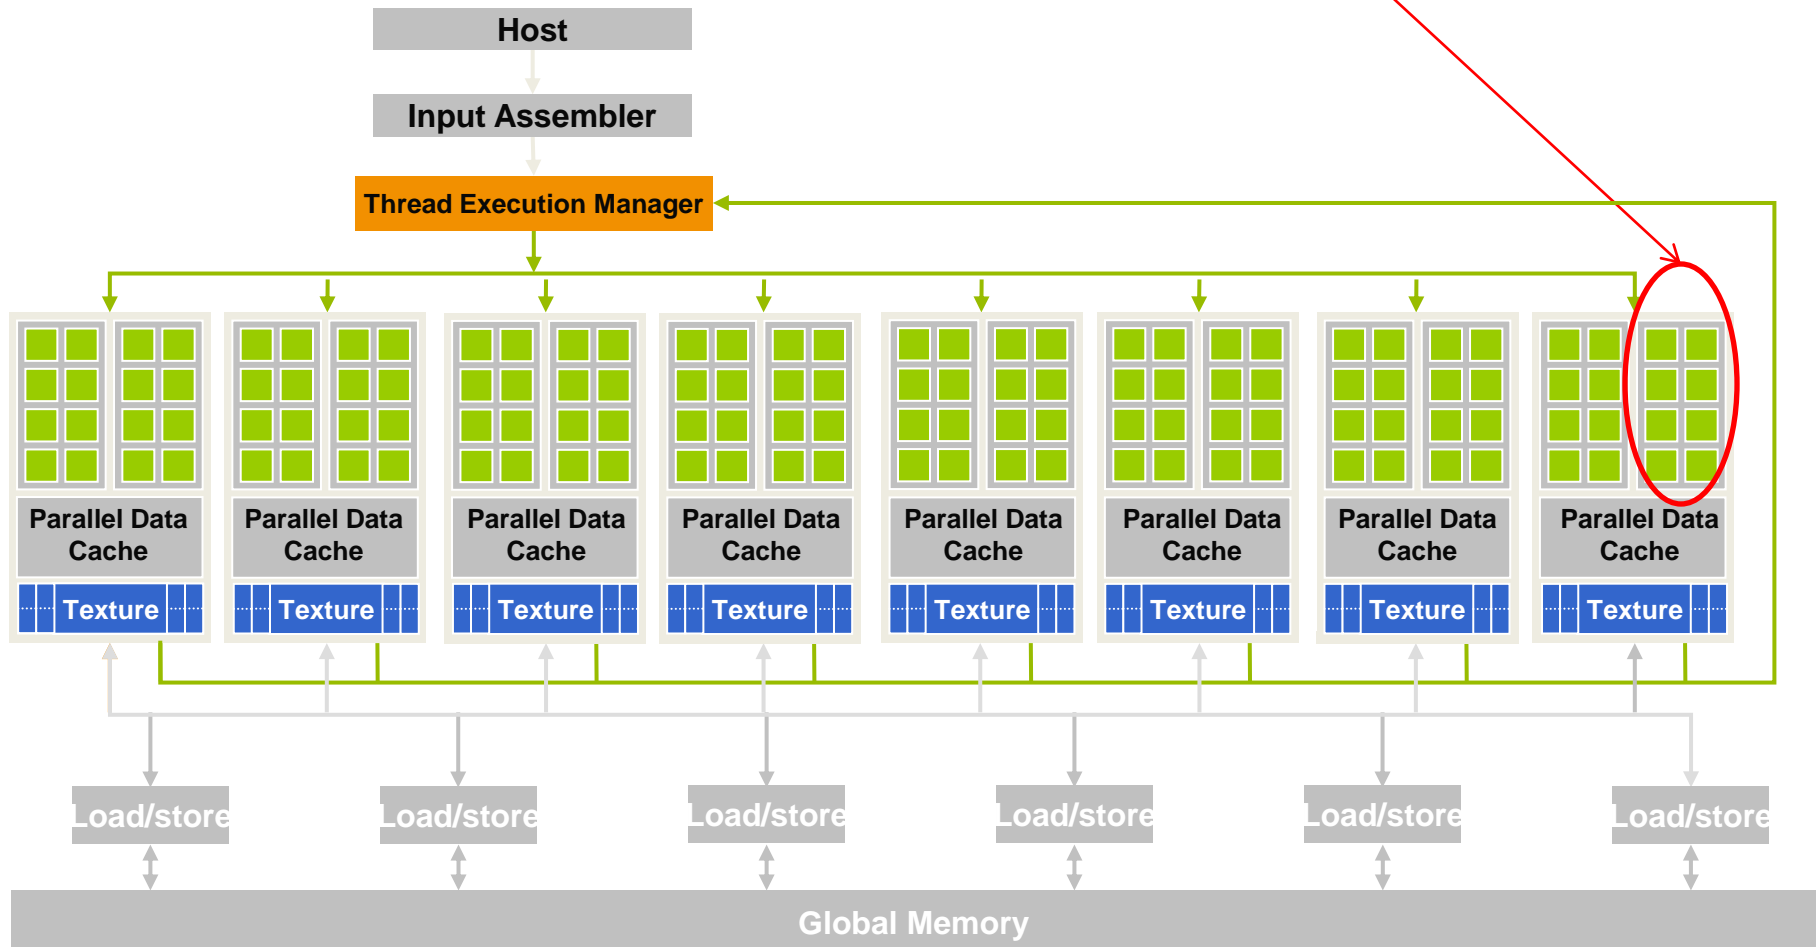
A Glimpse at A Modern GPU: GeForce 8800 (2007)

16 highly threaded SM's, >128 FPU's,
367 GFLOPS, 768 MB DRAM,
86.4 GB/S Mem BW,
4GB/S BW to CPU



A Glimpse at A Modern GPU

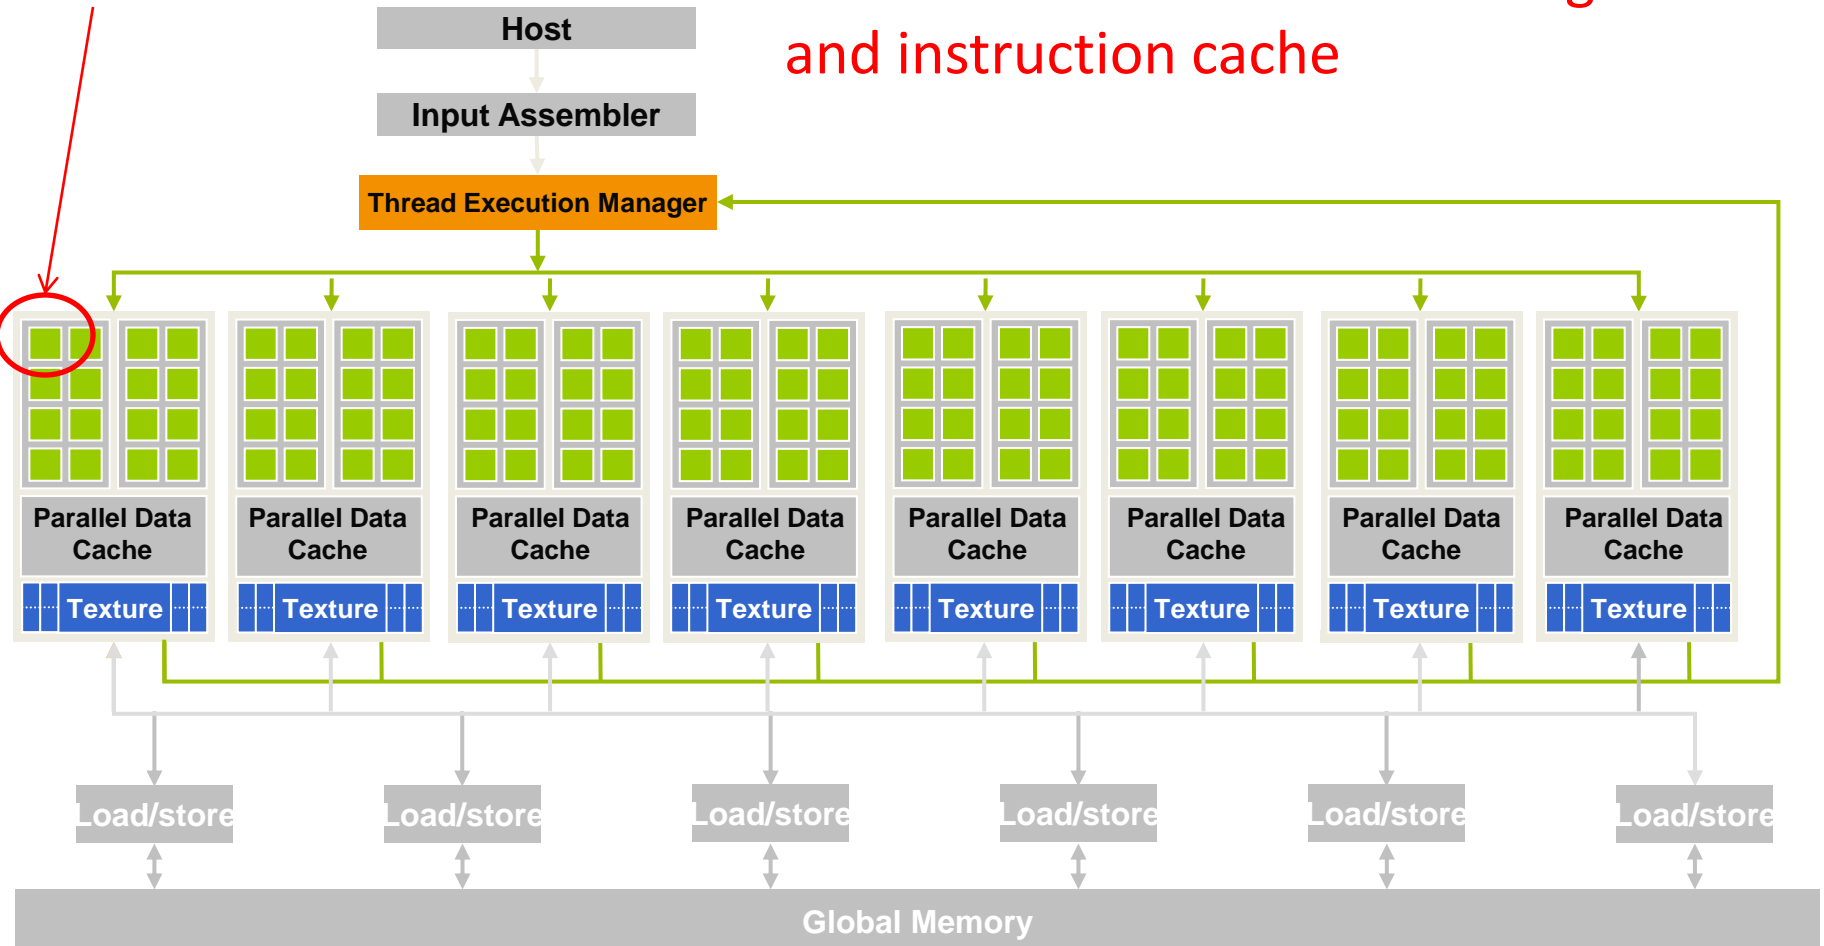
Streaming Multiprocessor (SM)



A Glimpse at A Modern GPU

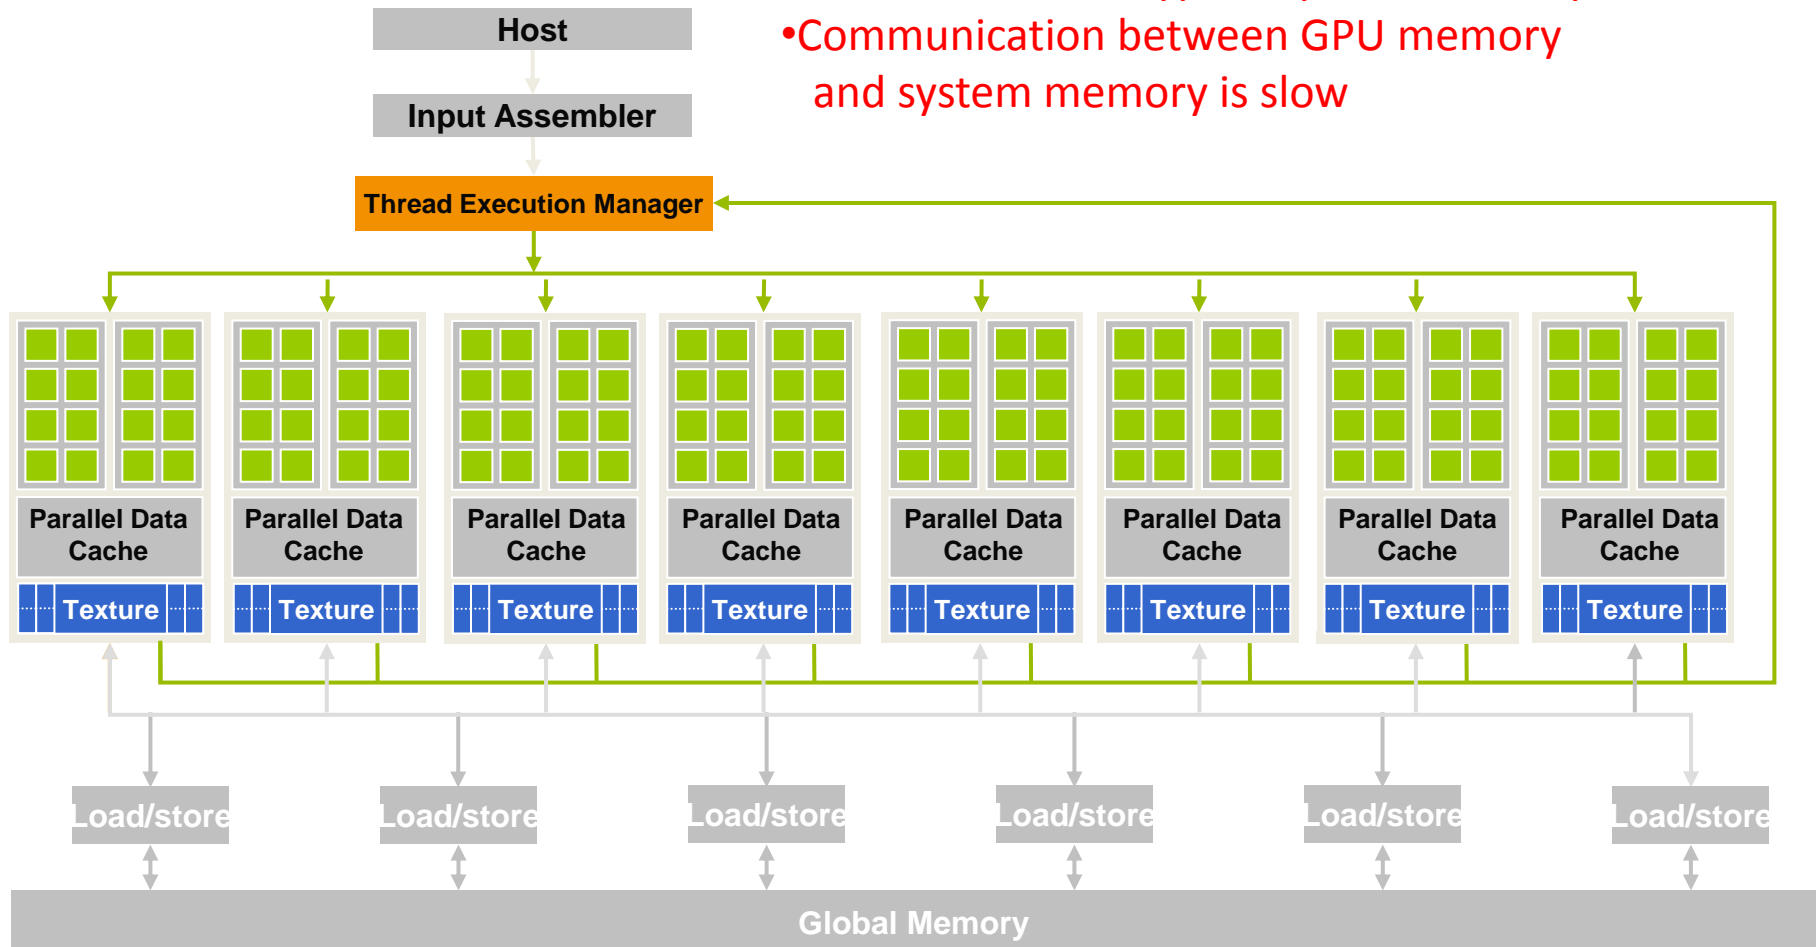
Streaming
Processor (SP)

SPs within SM share control logic
and instruction cache



A Glimpse at A Modern GPU

- Much higher bandwidth than typical system memory
- A bit slower than typical system memory
- Communication between GPU memory and system memory is slow



Amdahl's Law

$$\text{Execution Time After Improvement} = \text{Execution Time Unaffected} + (\text{Execution Time Affected} / \text{Amount of Improvement})$$

- Example:

"Suppose a program runs in 100 seconds on a machine, with multiply responsible for 80 seconds of this time. How much do we have to improve the speed of multiplication if we want the program to run 4 times faster?"

How about making it 5 times faster?

Improvement in your application speed depends on the portion that is parallelized

Two Main Things to Keep in Mind

- Try to increase the portion of your program that can be parallelized
- Figure out how to get around limited bandwidth of system memory

Enough for Today

- We are done with Chapter 1
- Some applications are better run on CPU while others on GPU
- The two main limitations
 - The parallelizable portion of the code
 - The communication overhead between CPU and GPU