



CSCI-GA.3033-012

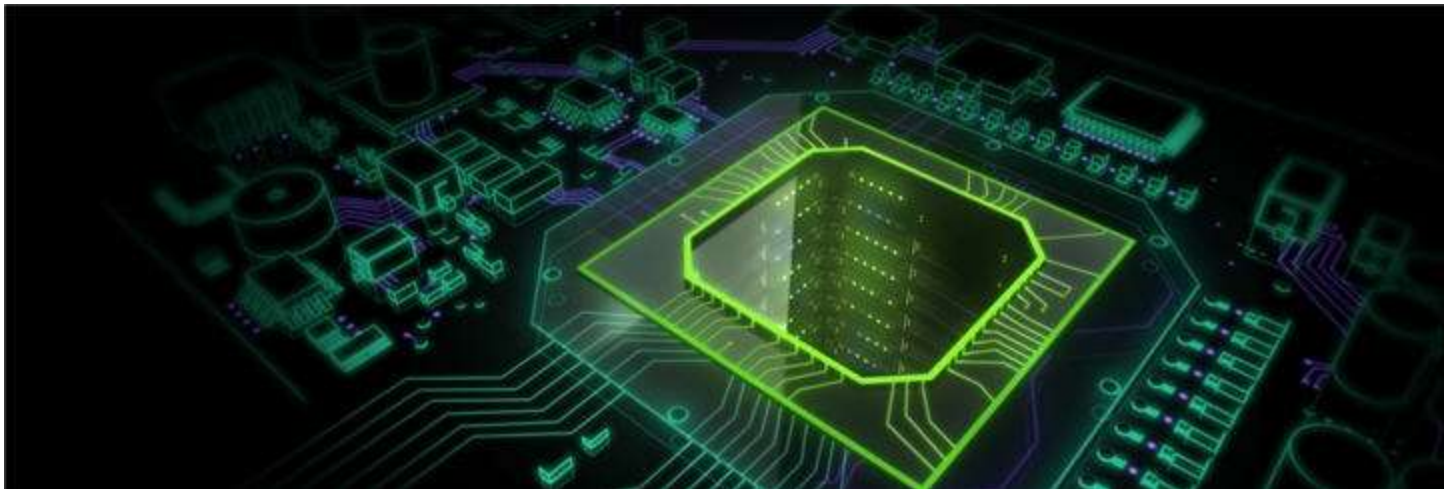
# Graphics Processing Units (GPUs): Architecture and Programming

## Lecture 13: Putting It All Together

Mohamed Zahran (aka Z)

[mzahran@cs.nyu.edu](mailto:mzahran@cs.nyu.edu)

<http://www.mzahran.com>




# This Lecture

We will pick a problem, **analyze** it, and see how it can be **written** and **optimized** for GPU.

# N-Body Problem

An N-body simulation numerically approximates the evolution of a system of bodies in which each body continuously interacts with every other body.



**Frames from an Interactive 3D Rendering of a 16,384-Body System**

# N-Body Problem

- Manifests itself in many domains: physics, astronomy, electromagnetics, molecules, etc.
- N points
- The answer at each point depends on data at all the other points
- $O(n^2)$
- To reduce complexity: compress data of groups of nearby points
  - A well-known algorithm to do this: Barnes Hut

# Barnes Hut n-Body Algorithm

Divided into 3 steps

1. Building the tree -  $O(n * \log n)$
2. Computing cell centers of mass -  $O(n)$
3. Computing Forces -  $O(n * \log n)$

# Barnes Hut n-Body Algorithm

0. Read input data and transfer to GPU

for each timestep do {

1. Compute bounding box around all bodies
2. Build hierarchical decomposition by inserting each body into octree
3. Summarize body information in each internal octree node
4. Approximately sort the bodies by spatial distance
5. Compute forces acting on each body with help of octree
6. Update body positions and velocities

}

7. Transfer result to CPU and output

**Executed on GPU**

# Barnes Hut n-Body Algorithm


0. Read input data and transfer to GPU

for each timestep do {

1. Compute bounding box around all bodies
2. Build hierarchical decomposition by inserting each body into octree
3. Summarize body information in each internal octree node
4. Approximately sort the bodies by spatial distance
5. Compute forces acting on each body with help of octree
6. Update body positions and velocities

}

7. Transfer result to CPU and output





# Barnes Hut n-Body Algorithm

0. Read input data and transfer to GPU

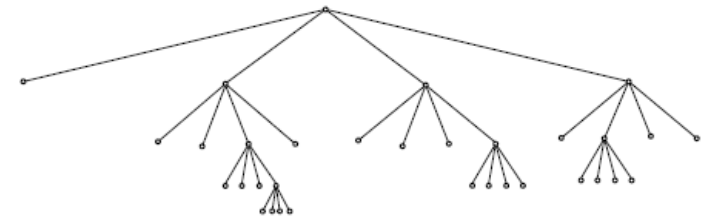
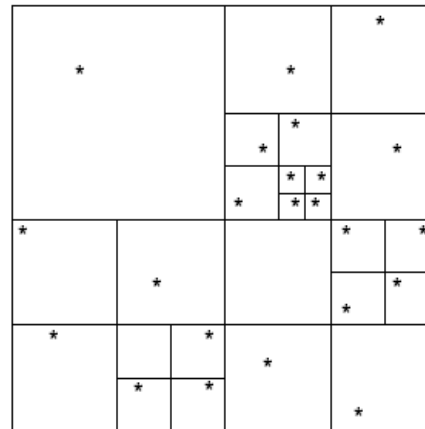
for each timestep do {

1. Compute bounding box around all bodies
2. Build hierarchical decomposition by inserting each body into octree
3. Summarize body information in each internal octree node
4. Approximately sort the bodies by spatial distance
5. Compute forces acting on each body with help of octree
6. Update body positions and velocities

}

7. Transfer result to CPU and output

**cells:** Internal tree nodes  
**Bodies:** leaves



# Barnes Hut n-Body Algorithm


0. Read input data and transfer to GPU

for each timestep do {

1. Compute bounding box around all bodies
2. Build hierarchical decomposition by inserting each body into octree
3. Summarize body information in each internal octree node
4. Approximately sort the bodies by spatial distance
5. Compute forces acting on each body with help of octree
6. Update body positions and velocities

}

7. Transfer result to CPU and output



# Barnes Hut n-Body Algorithm


0. Read input data and transfer to GPU

for each timestep do {

1. Compute bounding box around all bodies
2. Build hierarchical decomposition by inserting each body into octree
3. Summarize body information in each internal octree node
4. Approximately sort the bodies by spatial distance
5. Compute forces acting on each body with help of octree
6. Update body positions and velocities

}

7. Transfer result to CPU and output



Kernel 4 is not needed for correctness but for optimization


# First Step: Data Structure

- Dynamic data structures like trees are usually built using heap objects.
- Is it the best way to go?
- Drawbacks:
  - Access to heap objects is slow
  - Very hard to coalesce objects with multiple fields

How do we deal with this?

# First Step: Data Structure


- Use an array-based data structure
- To be able to coalesce:
  - use several aligned scalar arrays, one per field
- Array indices instead of pointers makes a faster code



# First Step: Data Structure

- Allocate bodies at the beginning and the cells at the end of the arrays
- Use an index of -1 as a "null pointer."
- Advantages.
  - A simple comparison of the array index with the number of bodies determines whether the index points to a cell or a body.
  - In some code sections, we need to find out whether an index refers to a body or to null. Because -1 is also smaller than the number of bodies, a single integer comparison suffices to test both conditions.
  - We can alias arrays that hold only cell information with arrays that hold only body information to reduce the memory consumption

# First Step: Data Structure



**b: body**

**c: cell**

**t: array length**

# Threads, Blocks, and Kernels

- The **thread count per block is maximized** and rounded down to the nearest multiple of the warp size for each kernel.
- All kernels **use at least as many blocks as there are streaming multiprocessors in the GPU**, which is automatically detected.
- Because all parameters passed to the kernels, such as the starting addresses of the various arrays, stay the same throughout the time step loop, we **copy them once into the GPU's constant memory**.
  - This is much faster than passing them with every kernel invocation.
- Data transferred from CPU to GPU only at the beginning of the program and at the end.
- code operates on octrees in which nodes can have up to eight children.
  - It contains many loops with a trip count of eight.
  - Loop unrolling is very handy here



# Kernel 1

- computes a bounding box around all bodies
  - The root of the octree
  - has to find the minimum and maximum coordinates in the three spatial dimensions
- Implementation:
  - break up the data into equal sized chunks and assigns one chunk to each block
  - Each block then performs a reduction operation
  - reduction is performed in shared memory in a way that avoids bank conflicts and minimizes thread divergence

# Kernel 2

- Implements an iterative tree-building algorithm that uses lightweight locks
- Bodies are assigned to the blocks and threads within a block in round-robin fashion.
- Each thread inserts its bodies one after the other by:
  - traversing the tree from the root to the desired last-level cell
  - attempting to lock the appropriate child pointer (an array index) by writing an otherwise unused value to it using an atomic operation
  - If the lock succeeds, the thread inserts the new body and release the lock

# Kernel 2

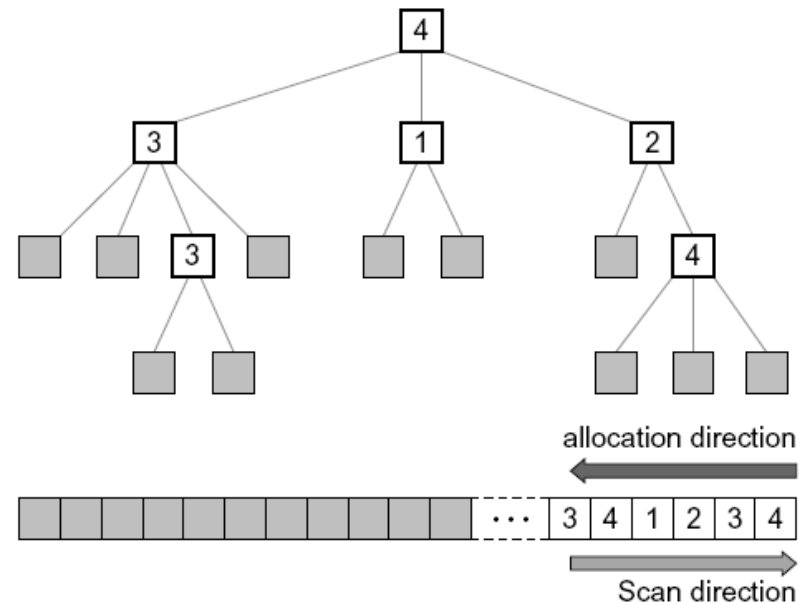
- If a body is already stored at this location, the thread:
  - creates a new cell by atomically requesting the next unused array index
  - inserts the original and the new body into this new cell
  - executes a memory fence ( *threadfence*) to *ensure the new subtree is visible to the rest of the cores*
  - attaches the new cell to the tree
  - releases the lock.

# Kernel 2

```
// initialize
cell = find_insertion_point(body); // nothing is locked, cell cached for retries
child = get_insertion_index(cell, body);
if (child != locked) {
    if (child == atomicCAS(&cell[child], child, lock)) {
        if (child == null) {
            cell[child] = body; // insert body and release lock
        } else {
            new_cell = ...; // atomically get the next unused cell
            // insert the existing and new body into new_cell
            __threadfence(); // make sure new_cell subtree is visible
            cell[child] = new_cell; // insert new_cell and release lock
        }
        success = true; // flag indicating that insertion succeeded
    }
}
__syncthreads(); // wait for other warps to finish insertion
```

# Kernel 3

- traverses the unbalanced octree from the bottom up to compute the center of gravity and the sum of the masses of each cell's children




# Kernel 3

- Cells are assigned to blocks and threads in a round-robin fashion.
  - Ensure load-balance
  - Start from leaves so avoid deadlocks
  - Allow some coalescing

# Kernel 5

- Requires the vast majority of the runtime
- For each body, the corresponding thread traverses some prefix of the octree to compute the force acting upon this body.



# Kernel 5

- Optimization: whenever a warp traverses part of the tree that some of the threads do not need, those threads are disabled due to thread divergence.
  - Make the union of the prefixes in a warp as small as possible
    - group spatially nearby bodies together
- Little computation to hide memory access
  - Optimization: Allow only one thread in a warp to read the pertinent data and cache them in shared memory.



# Summary of Optimizations

## MAIN MEMORY

### **Minimize Accesses**

- Let one thread read common data and distribute data to other threads via shared memory
- When waiting for multiple data items to be computed, record which items are ready and only poll the missing items
- Cache data in registers or shared memory
- Use thread throttling (see control-flow section)

### **Maximize Coalescing**

- Use multiple aligned arrays, one per field, instead of arrays of structs or structs on heap
- Use a good allocation order for data items in arrays

### **Reduce Data Size**

- Share arrays or elements that are known not to be used at the same time

### **Minimize CPU/GPU Data Transfer**

- Keep data on GPU between kernel calls
- Pass kernel parameters through constant memory

# Summary of Optimizations

## CONTROL FLOW

### Minimize Thread Divergence

- Group similar work together in the same warp

### Combine Operations

- Perform as much work as possible per traversal, i.e., fuse similar traversals

### Throttle Threads

- Insert barriers to prevent threads from executing likely useless work

### Minimize Control Flow

- Use compiler pragma to unroll loops

## LOCKING

### Minimize Locks

- Lock as little as possible (e.g., only a child pointer instead of entire node, only last node instead of entire path to node)

### Use Lightweight Locks

- Use flags (barrier/store and load) where possible
- Use atomic operation to lock but barrier/store or just store to unlock

### Reuse Fields

- Use existing data field instead of separate lock field

# Summary of Optimizations

## HARDWARE

### **Avoid Bank Conflicts**

- Control the accesses to shared memory to avoid bank conflicts

### **Use All Multiprocessors**

- Parallelize code across blocks
- Make the block count at least as large as the number of streaming multiprocessors

### **Maximize Thread Count**

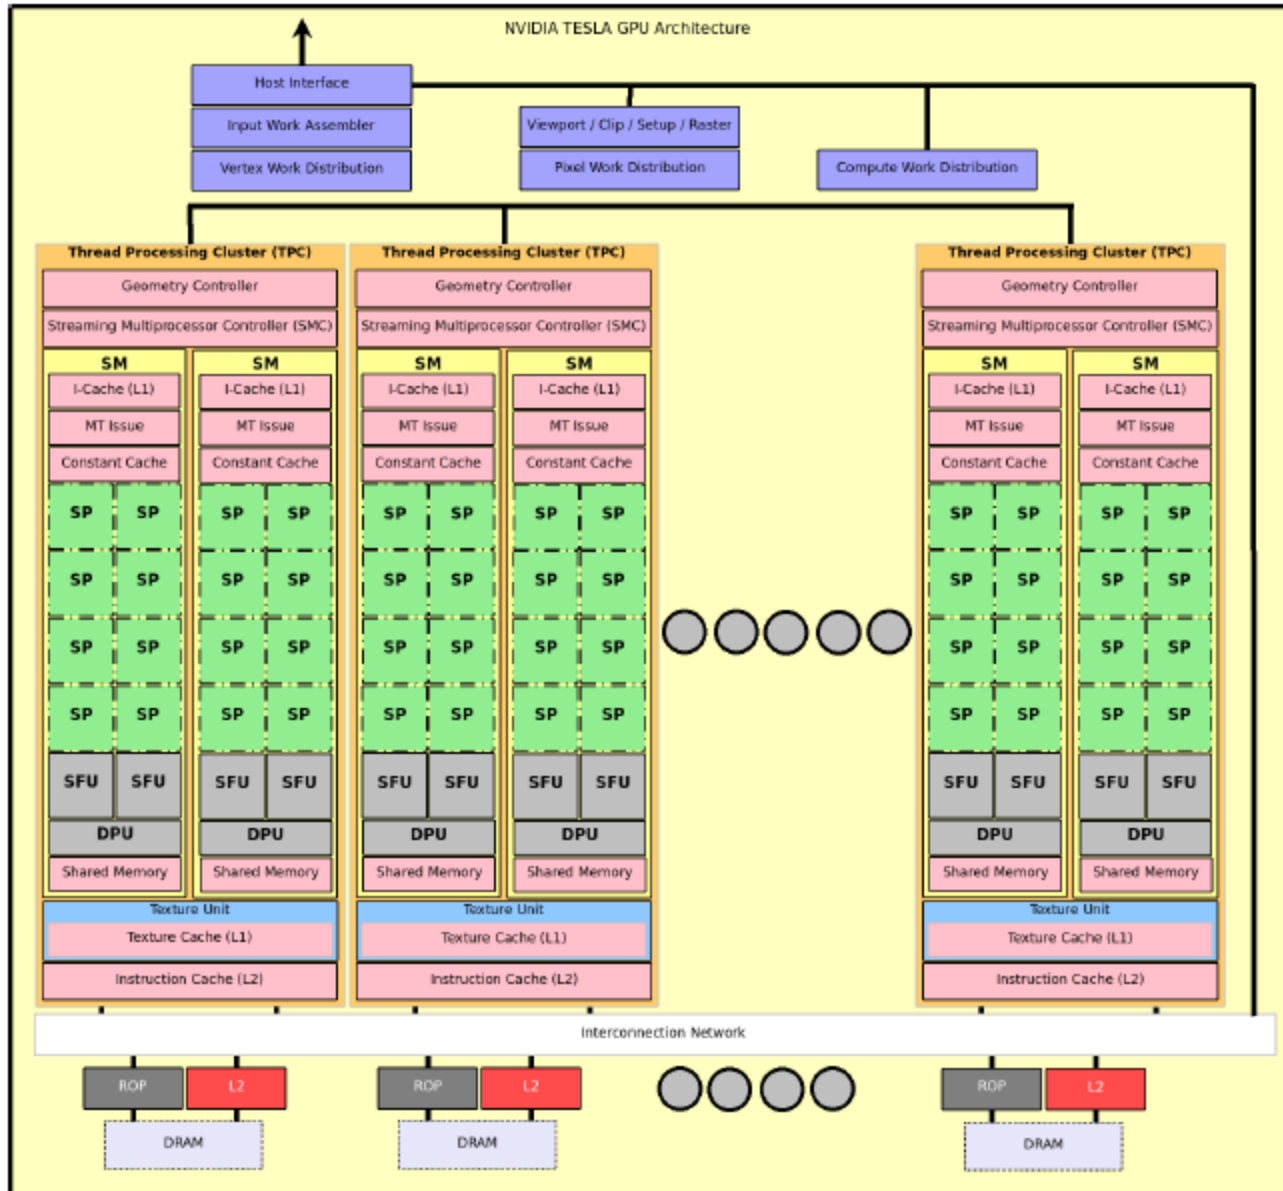
- Parallelize code across threads
- Limit shared memory and register usage to maximize thread count

The Road Ahead


# Reliability Issues

- Reliability in GPUs is not as addressed as other aspects
- Graphics applications may be fault tolerant, but other applications running on GPUs are not.
- Large scale GPU failure after shipment/deployment is not uncommon

# Typical GPU: Massive Parallelism

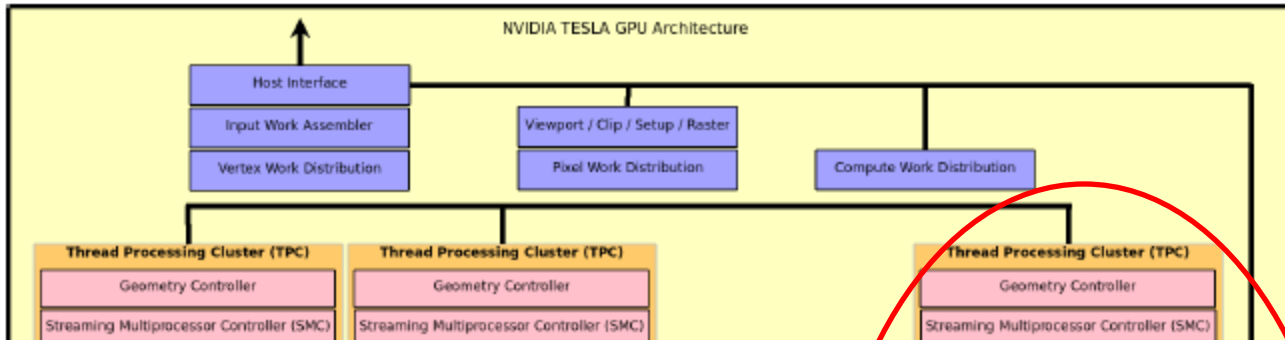


As the number of processing elements increases

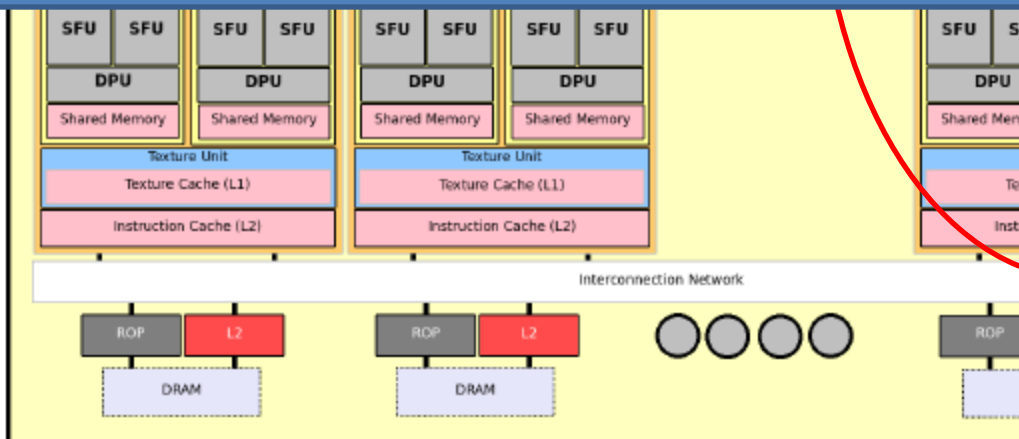


The probability that one or more of them fail increases.


# Typical GPU: Massive Parallelism

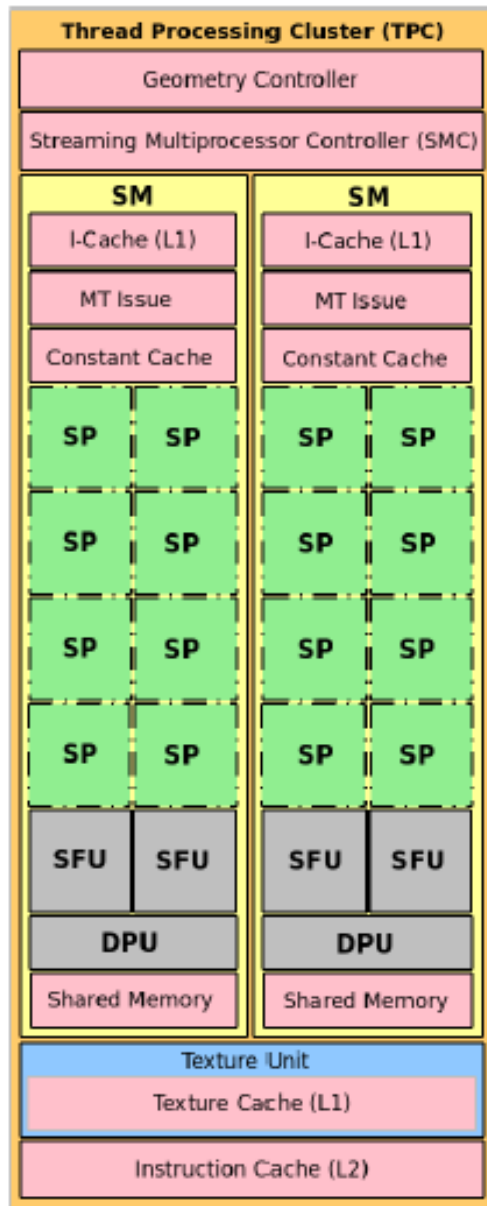


A loss of an SM in an 8-SM GPU can cause performance loss as high as 16%!



1 Faulty SM - Performance Penalty






## What if one of the SMs fails??

- Turn off the faulty part.
- Giving hints
- Reliability-aware software??



# GPUs in Embedded Devices

**Mobile SoC Market Share 2011**



# The Constraints of Mobile

- Energy
  - Cell phone battery capacity of 5-7 Watt-hour (tablets 20-40 Wh)
- Area
  - PCB size constraints
  - Cooling constraints

# Some Energy Numbers

Power Consumption Comparison		
	Apple iPhone 4 (AT&T)	Apple iPhone 4S (AT&T)
Idle	0.7W	0.7W
Launch Safari	0.9W	0.9W
Load AnandTech.com	<b>1.0W</b>	1.1W
Maps (Determine Current Location via GPS/WiFi)	<b>1.3W</b>	1.4W

Power Consumption Comparison		
	Apple iPhone 4 (AT&T)	Apple iPhone 4S (AT&T)
Launch Infinity Blade	<b>2.2W</b>	2.6W
Infinity Blade (Opening Scene, Steady State)	<b>2.0W</b>	2.2W

# Some Theoretical Performance Numbers

	Apple iPad 2	ASUS Transformer Prime	Some Nice Desktop
CPU	A5 @ 1GHz	Tegra 3 @ 1.4GHz	Sandy Bridge @ 3.4GHz
GPU	POWERVR SGX543MP2 @ 250MHz	Mobile GeForce @ 500MHz	GTX680 @ 1GHz
Memory Interface	64-bit @ (maybe) 800MHz = 6.4GB/s	32-bit	256-bit @ 6GHz = 192GB/s
GPU GFLOPS	16 GFLOPS	12 GFLOPS	3 TFLOPS


# Tegra 2 Mobile GeForce

- Separate vertex and pixel shaders
  - 4 of each, each capable of 1 multiply-add /clock
- Pixel, texture, vertex, and attribute caches
  - Reduce memory transactions
  - Pixel cache useful for UI components
- Memory controller optimizations
  - Arbitrate between CPU & GPU requests
  - Reorder requests to limit bank switching


# Case Study: the new iPad

- Screen resolution of 2048x1536
  - Higher than nearly all desktop and laptop displays
- Battery life approximately equal to previous version
- Resolution, power, performance: pick any two.

# iPad Gaming Performance







Apple A5X Die Shot

Image from [UBMTechInsights](http://UBMTechInsights.com)



# Interesting Question

Do we need GPU for non-graphics application in mobile gadgets?

# Conclusions

- When considering your problem:
  - Pick your algorithm
  - Choose the data structure
  - Try to make as many threads and blocks as possible busy
  - Know your hardware
  - Tweaks are inevitable
  - Correctness, performance, and power.